

# Linear Regression - Student Percentage Prediction

```
In [143... # Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [144... df = pd.read_csv('student_scores.csv')
df.head()
```

```
Out[144... 
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

## Data Splitting

```
In [145... df.shape
```

```
Out[145... (25, 2)
```

```
In [146... X,y = df['Hours'],df['Scores']
```

```
In [147... from sklearn.model_selection import train_test_split
```

```
In [148... X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

# Exploratory Data Analysis

```
In [149... X_train.shape
```

```
Out[149... (20,)
```

```
In [150... X_train.dtypes
```

```
Out[150... dtype('float64')
```

```
In [151... X_train.info()
```

```
<class 'pandas.core.series.Series'>  
Index: 20 entries, 9 to 6  
Series name: Hours  
Non-Null Count  Dtype  
-----  
20 non-null     float64  
dtypes: float64(1)  
memory usage: 320.0 bytes
```

```
In [152... X_train.describe()
```

```
Out[152... count    20.000000  
mean      4.960000  
std       2.566874  
min       1.100000  
25%      3.075000  
50%      4.650000  
75%      7.475000  
max      9.200000  
Name: Hours, dtype: float64
```

```
In [153... X_train.isnull().sum()
```

```
Out[153... 0
```

```
In [154... X_train[X_train.duplicated()]
```

```
Out[154... 20    2.7  
           Name: Hours, dtype: float64
```

```
In [159... y_train.drop(labels=[20],inplace=True)
```

```
In [160... X_train.drop_duplicates(inplace=True)
```

```
In [161... X_train.shape
```

```
Out[161... (19,)
```

```
In [162... lower,upper = np.percentile(X_train,25),np.percentile(X_train,75)  
iqr = upper-lower  
lb,ub = iqr-1.5*lower,iqr+1.5*upper
```

```
In [163... lower,upper,iqr
```

```
Out[163... (3.25, 7.550000000000001, 4.300000000000001)
```

```
In [164... lb,ub
```

```
Out[164... (-0.5749999999999993, 15.625000000000002)
```

```
In [165... X_train[X_train<lb],X_train[X_train>ub]
```


```
Out[165... (Series([], Name: Hours, dtype: float64),  
          Series([], Name: Hours, dtype: float64))
```

- No Outliers Found

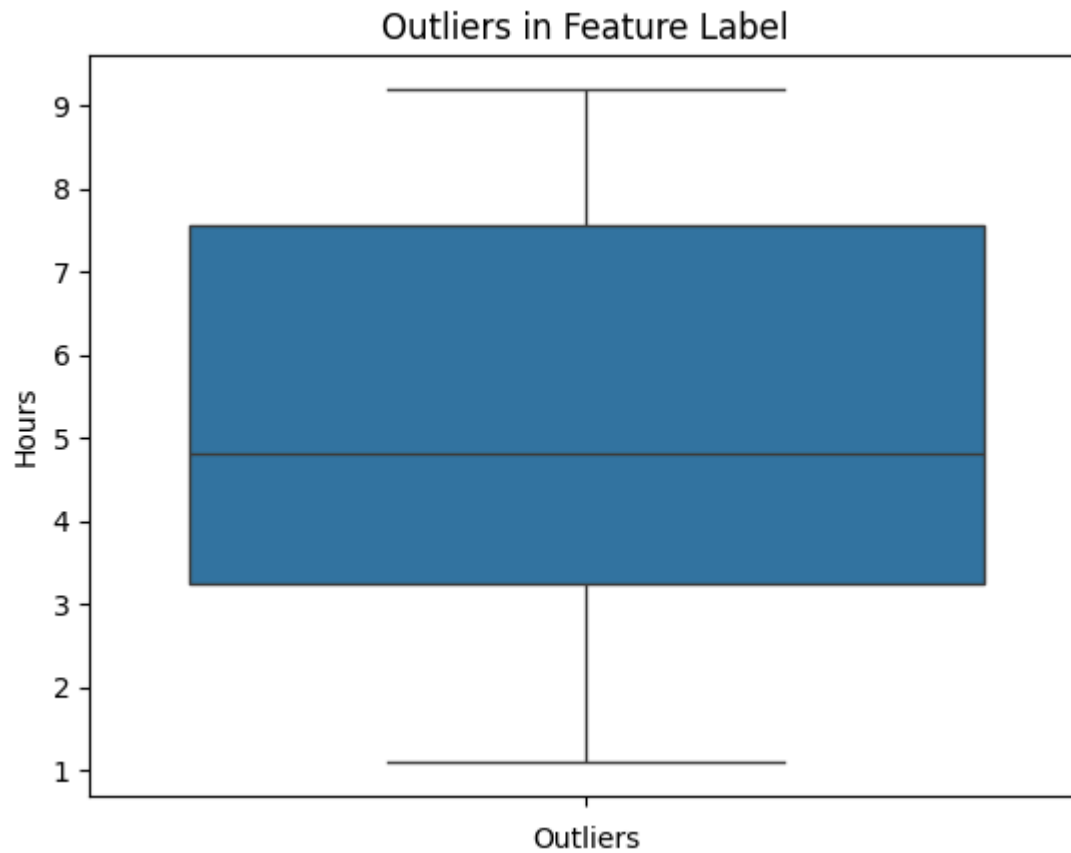
```
In [166... X_train.head()
```

```
Out[166... 9      2.7
          13     3.3
           1     5.1
          22     3.8
           5     1.5
          Name: Hours, dtype: float64
```

```
In [167... sns.histplot(X_train,bins=8,kde=True)
          plt.xlabel('Hours')
          plt.ylabel('Counts')
          plt.title('Hours Distribution')
          plt.tight_layout()
          plt.grid()
          plt.show()
```



```
In [168... sns.boxplot(data=X_train)
plt.xlabel('Outliers')
plt.ylabel('Hours')
plt.title('Outliers in Feature Label')
plt.show()
```



- No Outliers

## Building Model

In [170... `# Normalization`

In [174... `X_train = pd.DataFrame(X_train)`

In [183... `X_test = pd.DataFrame(X_test)`

```
In [184... from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train_normalized = scaler.fit_transform(X_train)
X_test_normalized = scaler.transform(X_test)
```

```
In [185... from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train_normalized, y_train)
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)
```

Coefficients: [78.59073018]  
Intercept: 13.341395700266219

```
In [210... fig, axes = plt.subplots(1, 2, figsize=(12, 5))

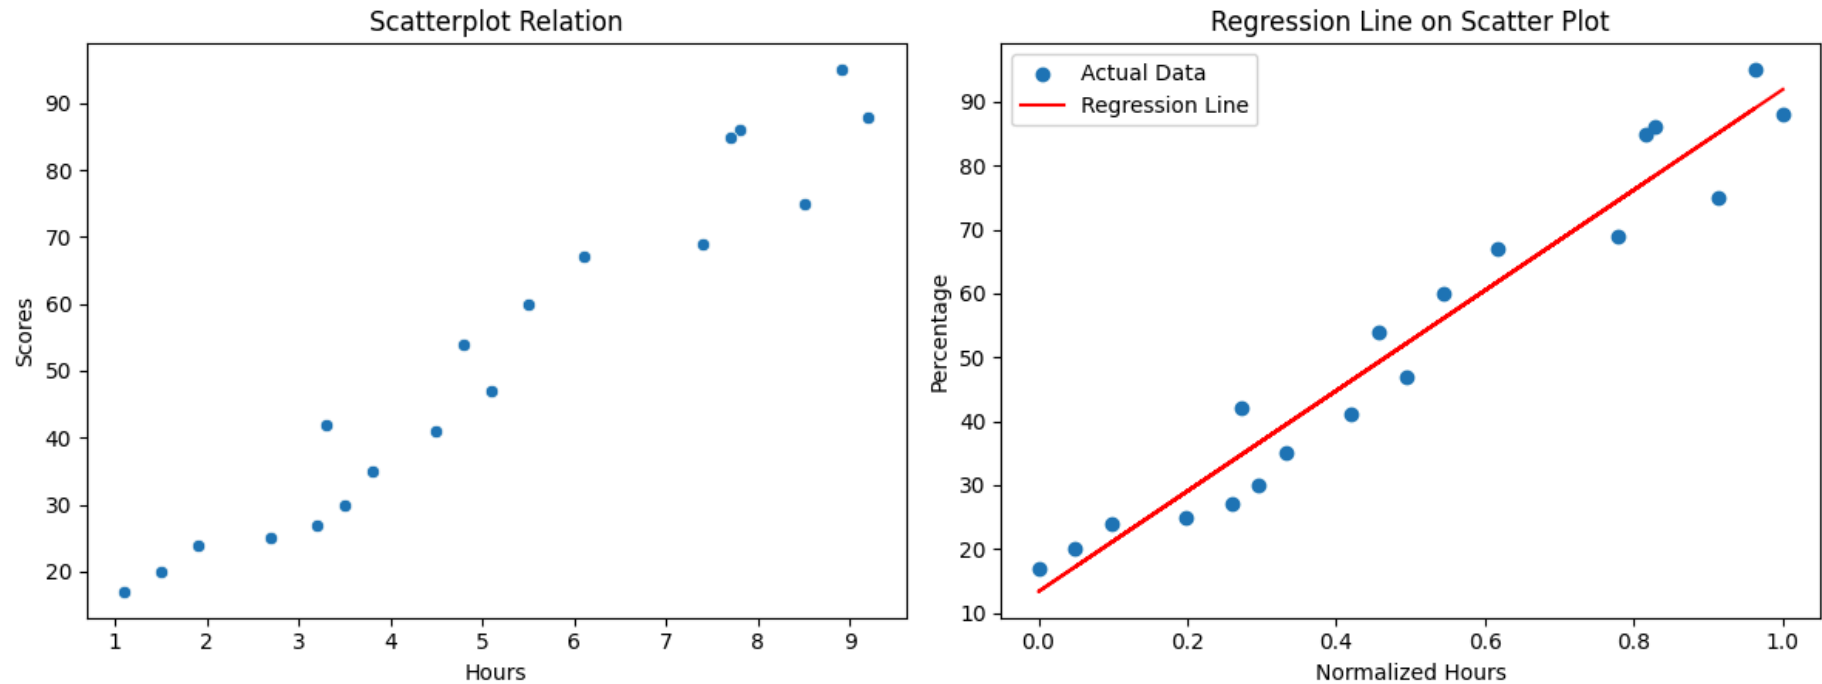
# 1st Scatter plot
sns.scatterplot(x=X_train['Hours'], y=y_train, ax=axes[0])
axes[0].set_xlabel('Hours')
axes[0].set_ylabel('Scores')
axes[0].set_title('Scatterplot Relation')

# 2nd Scatter plot with Regression Line
axes[1].scatter(X_train_normalized[:, 0], y_train, label='Actual Data')

axes[1].plot(X_train_normalized[:, 0], model.predict(X_train_normalized), color='red', label='Regression Line')

# Add Labels and title
axes[1].set_xlabel('Normalized Hours')
axes[1].set_ylabel('Percentage')
axes[1].set_title('Regression Line on Scatter Plot')
axes[1].legend()
plt.suptitle('Scatterplots Along with Regression Line')
plt.tight_layout()
plt.show()
```

## Scatterplots Along with Regression Line



## Evaluation

```
In [211... from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
y_pred = model.predict(X_test_normalized)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)
print("Root Mean Squared Error (RMSE):", rmse)
print("R-squared (R2):", r2)
```



Mean Absolute Error (MAE): 3.933980548404949  
Mean Squared Error (MSE): 18.901095098644035  
Root Mean Squared Error (RMSE): 4.347538970342191  
R-squared (R2): 0.9678771327351393

## Prediction

In [212...

```
import warnings
warnings.filterwarnings("ignore")
hours_to_predict = float(input('Enter Hours to predict Percentage..'))
normalized_hours_to_predict = scaler.transform([[hours_to_predict]])

predicted_percentage = model.predict(normalized_hours_to_predict)

print(f"Predicted Percentage for {hours_to_predict} hours / day of study is : {round(predicted_percentage[0],2)}")
```

Predicted Percentage for 5.0 hours / day of study is : 51.18